CSYE7374 BigData Systems and Intelligent Analytics
Team 4 - Assignment 2

Steps for creating EMR cluster from Amazon Command Line Interface (AWS CLI)

Note: Since we are using Apple Macbook, we have provided the OS X version. Linux supports the same commands.

Type the following commands in Terminal in the same sequence.

1. $ sudo pip install awscli

-- The below command will ask for inputs

2. $ aws configure

    a. AWS Access Key ID [None]: Put in your key
    b. AWS Secret Access Key [None]: Put in your key
    c. Default region name [None]: us-east-1
    d. Default output format [None]: json

-- The below command will create a cluster with the name 'Development Cluster'

3. $ aws emr create-cluster --release-label=emr-4.0.0 --instance-type=m3.xlarge --instance-count=1 --applications Name=Spark Name=Hadoop --ec2-attributes KeyName=YourKey

-- This command will list all your clusters, starting from the most recent one at top.

4. $ aws emr list-clusters

Copy the cluster ID from the description

-- This command will give all the details about your cluster, including its status and SSH details.

5. $ aws emr describe-cluster --cluster-id yourClusterID

Copy Public Master DNS to do SSH

6. $ ssh hadoop@copiedpublicdns

- If everything went fine, you should be all set and into the EMR Master node.

1. **Steps for setting up Kafka on EMR Cluster Master Node (Scala Example)**

- Prerequisite: You need to open at least **7** terminal windows connected to EMR
- Note- Kafka doesn't work on Spark 1.4.1 so we installed Spark 1.4.0

- EMR Master node doesn't have **wget** installed. so first we'll have to install that in order to be able to download any other stuff
    - $ sudo yum -y install wgetv
    - $ wget --no-check-certificate --no-cookies --header "Cookie: oraclelicense=accept-securebackup-cookie"http://download.oracle.com/otn-pub/java/jdk/7u67-b01/jdk-7u67-linux-x64.rpm

- Now we need to download Spark 1.4.0 as EMR comes with Spark 1.4.1
    - $ wget http://mirror.symnds.com/software/Apache/spark/spark-1.4.0/spark-1.4.0-bin-hadoop2.6.tgz
    - $ tar -xzf spark-1.4.0-bin-hadoop2.6.tgz

- Then we can download Kafka 2.11-0.8.2.1 from the mirror site using wget
    - $ wget http://apache.mirrors.tds.net/kafka/0.8.2.1/kafka_2.9.1-0.8.2.1.tgz
    - $ tar -xzf kafka_2.9.1-0.8.2.1.tgz
    - $ cd kafka_2.9.1-0.8.2.1

- Now we will start Zookeeper (You have to be in Kafka's installation directory)
    - $ bin/zookeeper-server-start.sh config/zookeeper.properties

- Next, in another terminal, we will start a Kafka Broker
    - $ cd kafka_2.9.1-0.8.2.1
    - $ bin/kakfa-server-start.sh config/server.properties

- In another terminal window, we will create a Kafka topic
    - $ cd kafka_2.9.1-0.8.2.1
    - $ bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1 --topic kafkatopicEMR1

- In another terminal window, we will start a producer
    - $ cd kafka_2.9.1-0.8.2.1

    - $ bin/kafka-console-producer.sh --broker-list localhost:9092 --topic kafkatopicEMR1

- Now we will start a consumer, in another terminal
    - $ cd kafka_2.9.1-0.8.2.1
    - $ bin/kafka-console-consumer.sh --zookeeper localhost:2181 --topic kafkatopicEMR1 --from-beginning

- Now in another terminal, lets navigate to Spark 1.4.0 directory and run a producer
    - $ cd usr/lib/spark-1.4.0/
    - $ bin/run-example org.apache.spark.examples.streaming.KafkaWordCountProducer localhost:9092 kafkatopicEMR1 10 5

- Now in another terminal, lets navigate to Spark 1.4.0 directory and run the wordcount problem
    - bin/run-example org.apache.spark.examples.streaming.KafkaWordCount localhost:2181 myconsumergroup kafkatopicEMR1 1

NOTE:
If firewall is activated

[hadoop@ip-10-238-177-148 ebs]$ sudo service iptables save
[hadoop@ip-10-238-177-148 ebs]$ sudo service iptables stop
[hadoop@ip-10-238-177-148 ebs]$ sudo chkconfig iptables off

2. Steps for executing Flume wordcount example (in Scala)

- Install Flume 1.6.0 by using the following commands:

    $ sudo wget http://www.gtlib.gatech.edu/pub/apache/flume/1.6.0/apache-flume-1.6.0-bin.tar.gz
    $ sudo tar -xzvf apache-flume-1.6.0-bin.tar.gz
    $ cd apache-flume-1.6.0-bin

- Configure the config file in the flume as follows:

    $ vi conf/flume-conf.properties.template

```
#Enter the following lines in the config file and save it.

a1.sources = r1
a1.sinks = k1
a1.channels = c1

#Describe/configure the source

a1.sources.r1.type = avro
a1.sources.r1.channels = c1
a1.sources.r1.bind = localhost

#Flume startup 1999,which wait for avro client to connect to it and
send Avro Flume event

a1.sources.r1.port = 1999

# Describe the sink

a1.sinks = k1
a1.sinks.k1.type = avro
a1.sinks.k1.hostname = localhost

###9999 is opened by other process, Flume will write data to it via
Socket

a1.sinks.k1.port = 9999

# Use a channel which buffers events in memory

a1.channels.c1.type = memory
a1.channels.c1.capacity = 1000
a1.channels.c1.transactionCapacity = 100

# Bind the source and sink to the channel

a1.sources.r1.channels = c1
a1.sinks.k1.channel = c1
```

- Start the agent by using the following command:

```
$ bin/flume-ng agent --conf conf --conf-file
conf/flume-conf.properties.template --name a1
```

- Start the receiver which creates a server and listens for flume events by the following command:

$ bin/run-example org.apache.spark.examples.streaming.FlumeEventCount localhost 9999

- Create a Client which acts as the source:

$ bin/flume-ng avro-client -F /home/ubuntu/spark-1.4.0-bin-hadoop2.6/examples/src/main/resources/people.jsonfileSuffix -H localhost -p 9999

The output will be something like this:

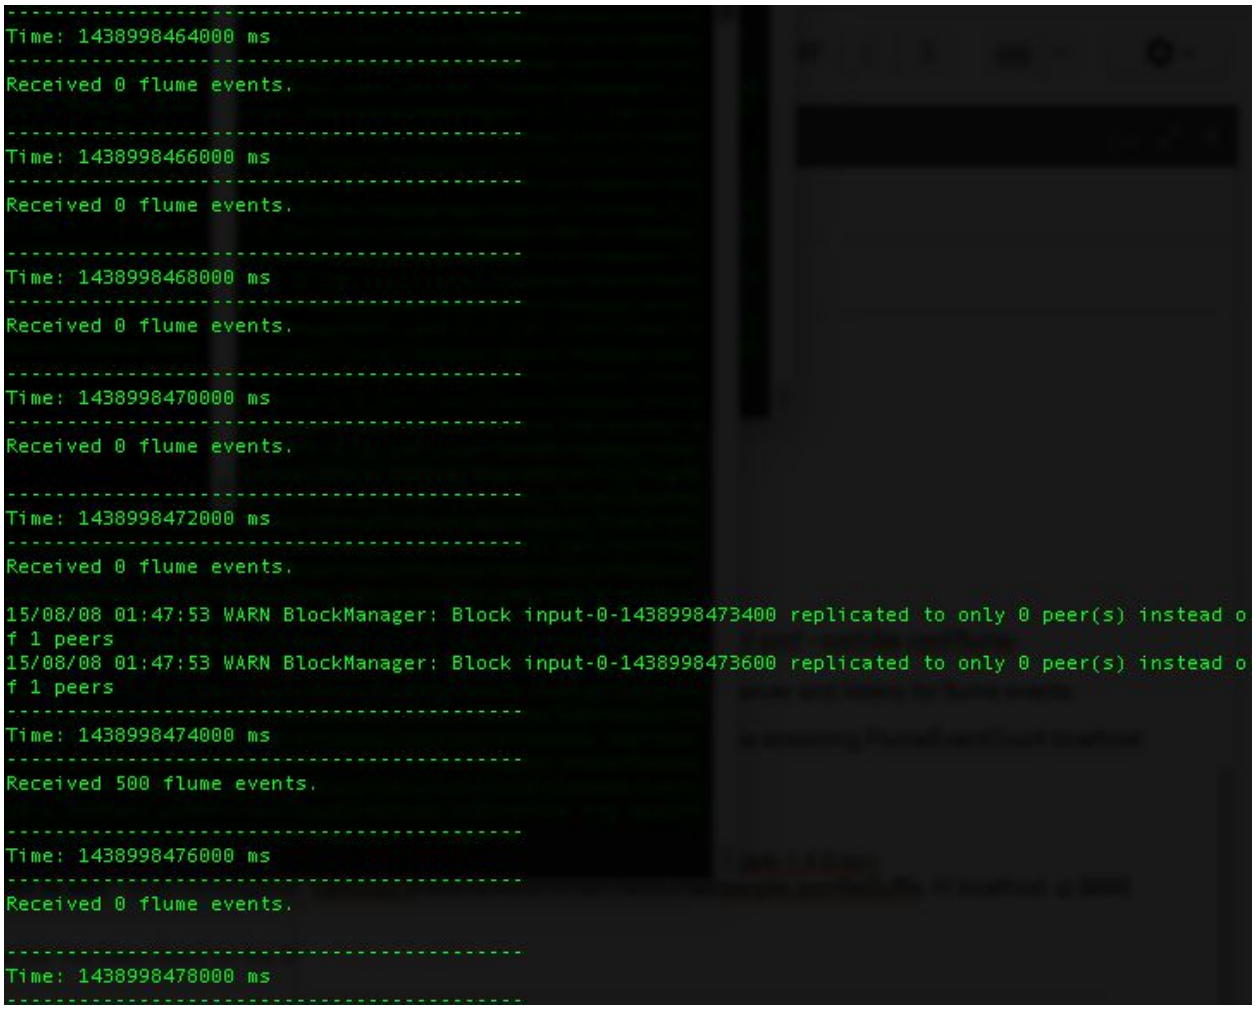

3. Steps for executing Kinesis wordcount example (in Scala)

- Note: Kinesis doesn't run on pre-built Spark so we had to build Spark 1.4.1 from source using maven. We tried doing this on EMR but building errored out with OutOfMemory exception. Hence, we created Kinesis stream on our local itself.

- Let's create a Kinesis Stream named Stream3
    - $ aws kinesis create-stream --stream-name Stream3 --shard-count 1

- Now let's take a look at our stream details
    - $ aws kinesis describe-stream --stream-name Stream3

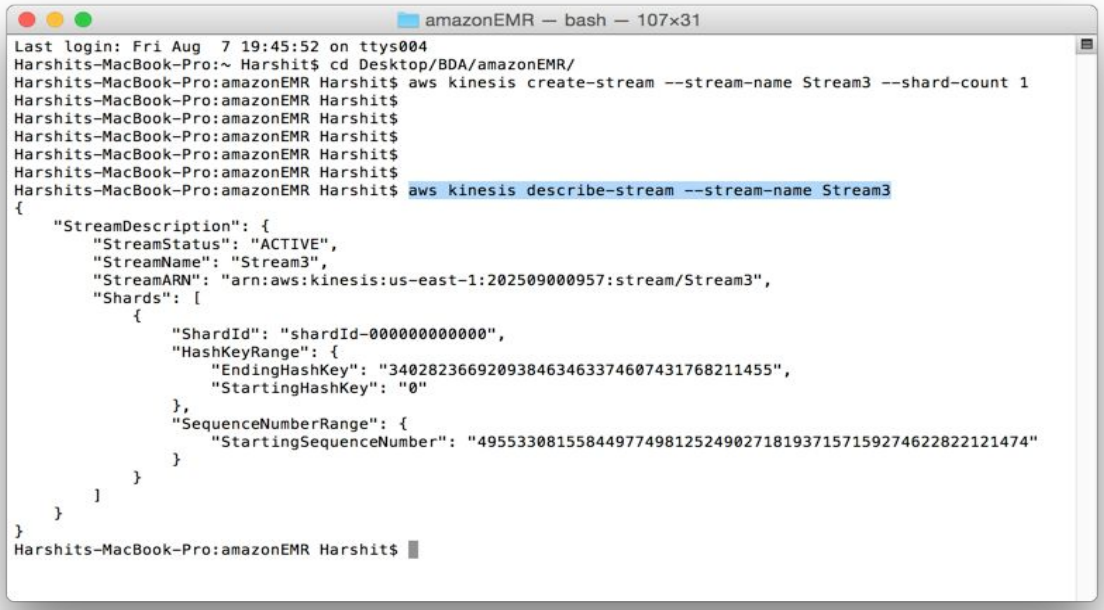

- Open a new terminal and let's run the WordCount example from Spark 1.4.1 directory (manually built version)
    - $ bin/run-example streaming.KinesisWordCountASL myapp3 Stream3 https://kinesis.us-east-1.amazonaws.com

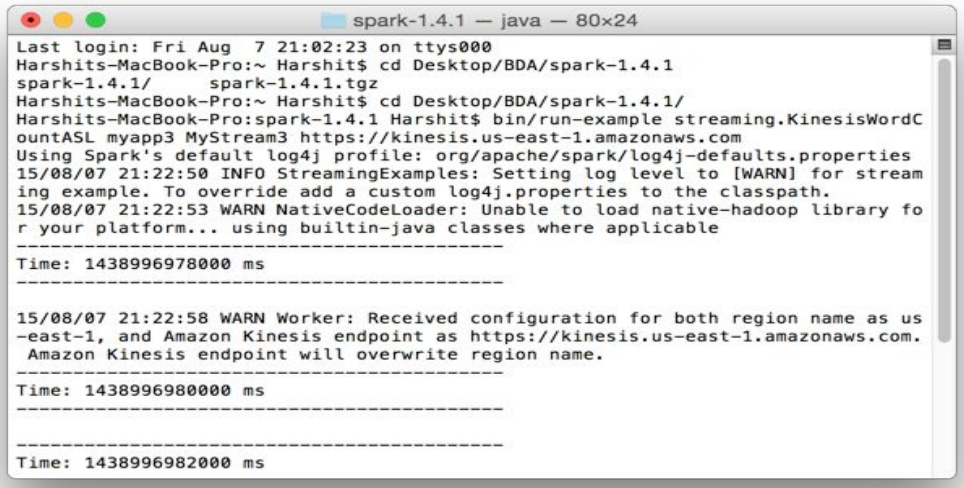

- Open another terminal and run the Kinesis Producer (Spark 1.4.1 directory)
    - $ bin/run-example streaming.KinesisWordProducerASL Stream3 https://kinesis.us-east-1.amazonaws.com 1000 10

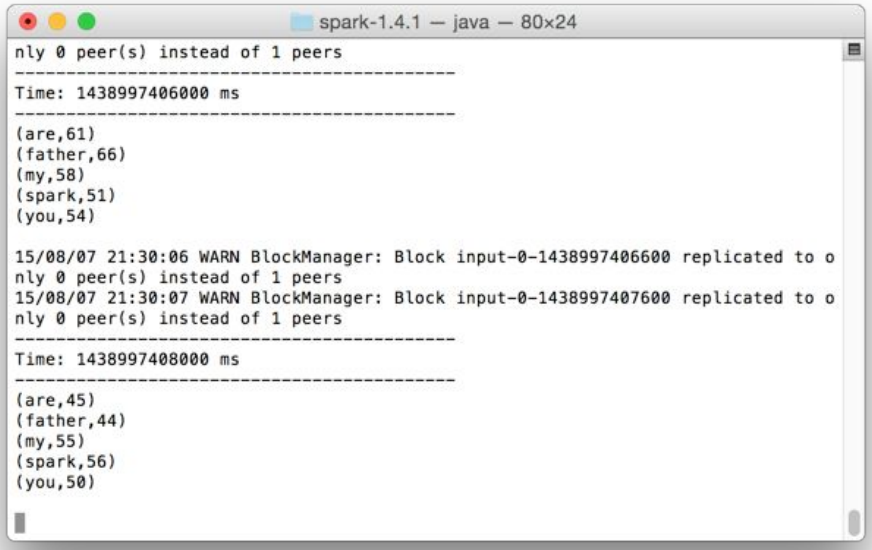

- After that, we can delete the stream
    - $ aws kinesis delete-stream --stream-name Stream3

4. Steps for executing HDFS wordcount example (in Scala)

- Change the directory to hadoop and create a folder called input in the hadoop folder as follows:

  $ cd /usr/lib/hadoop
  $ bin/hadoop fs -mkdir /input

- Run the example from the spark directory with the hdfs directory as input:

  $ bin/run-example org.apache.spark.examples.streaming.HdfsWordCount hdfs:///input

- Put a file into the hdfs directory for wordcount:

  $ bin/hadoop fs -put /usr/lib/spark/examples/src/main/resources/kv1.txt /input

- The output would be something like this:

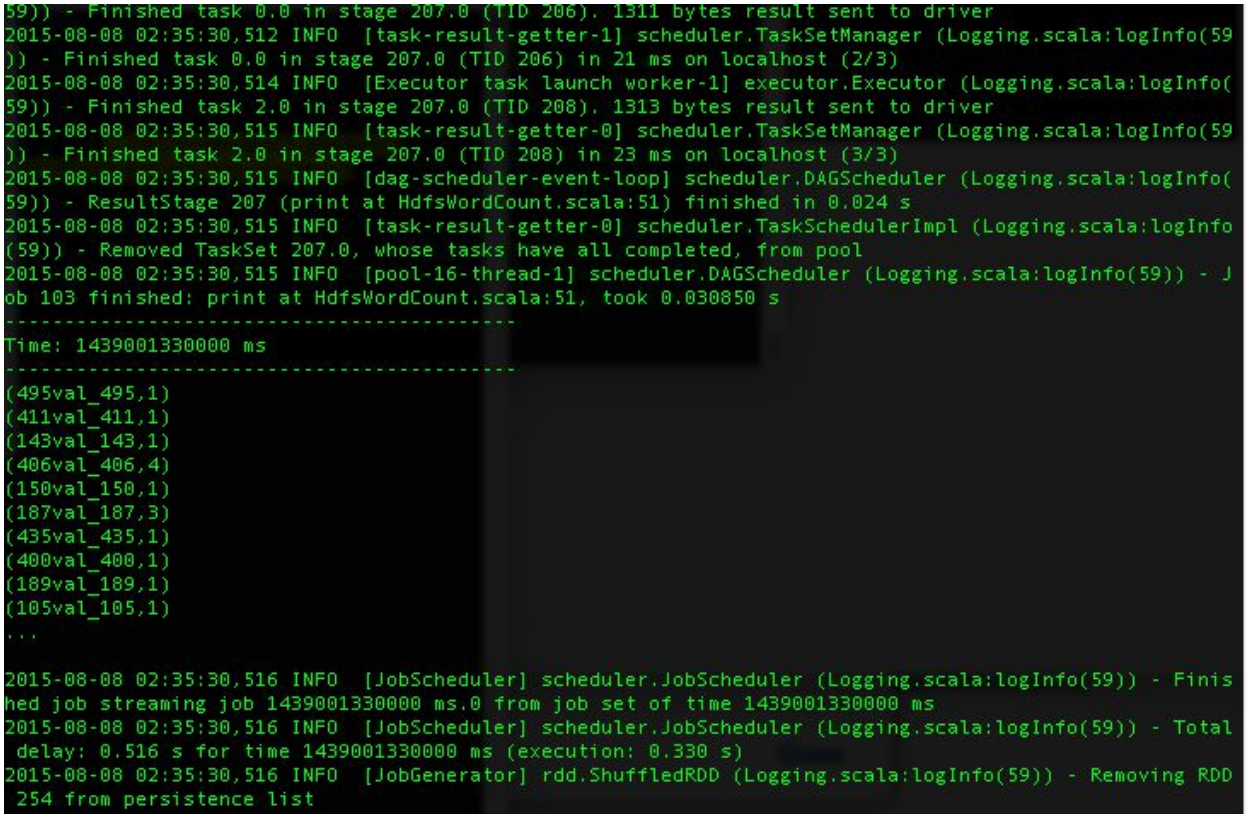

5. Steps for executing MQTT Hello world example (in Scala)

- Install MQTT by using the following command:

  $ apt-get install mosquitto

- Run the publisher as follows:

$ bin/run-example org.apache.spark.examples.streaming.MQTTPublisher tcp://localhost:1883 foo

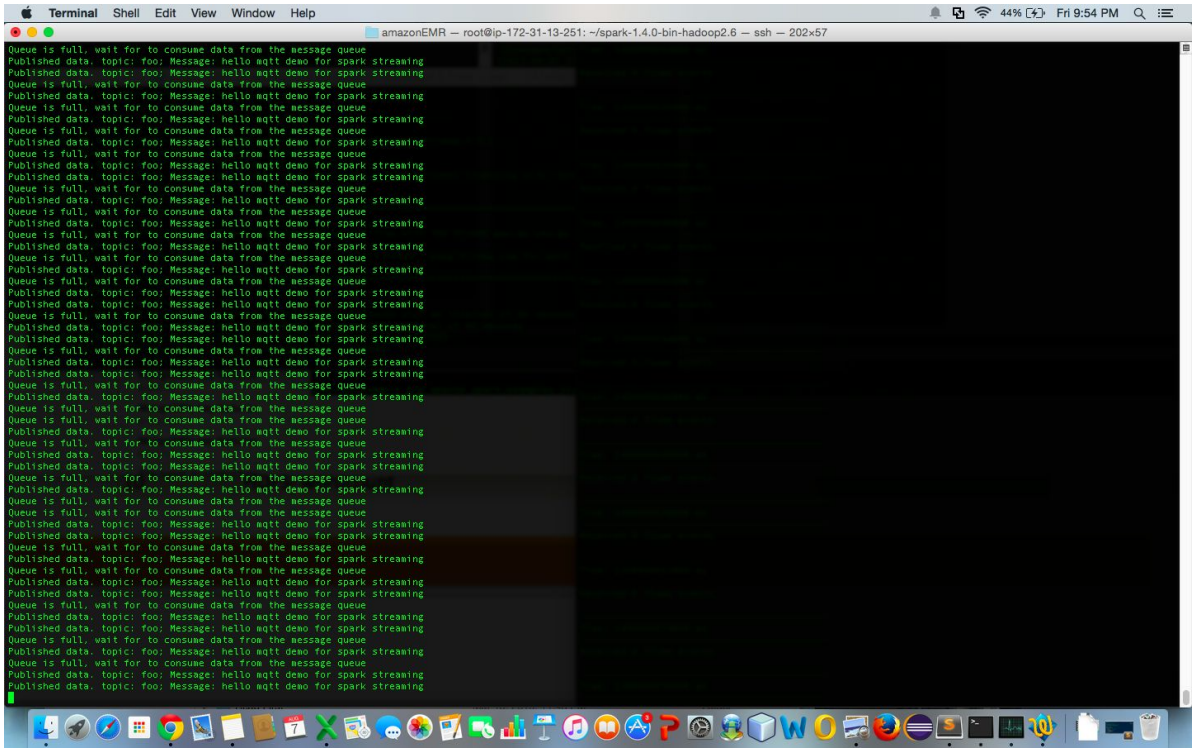

- Run the example as follows:

  $ bin/run-example org.apache.spark.examples.streaming.MQTTWordCount tcp://localhost:1883 foo

- The output will be something like this:

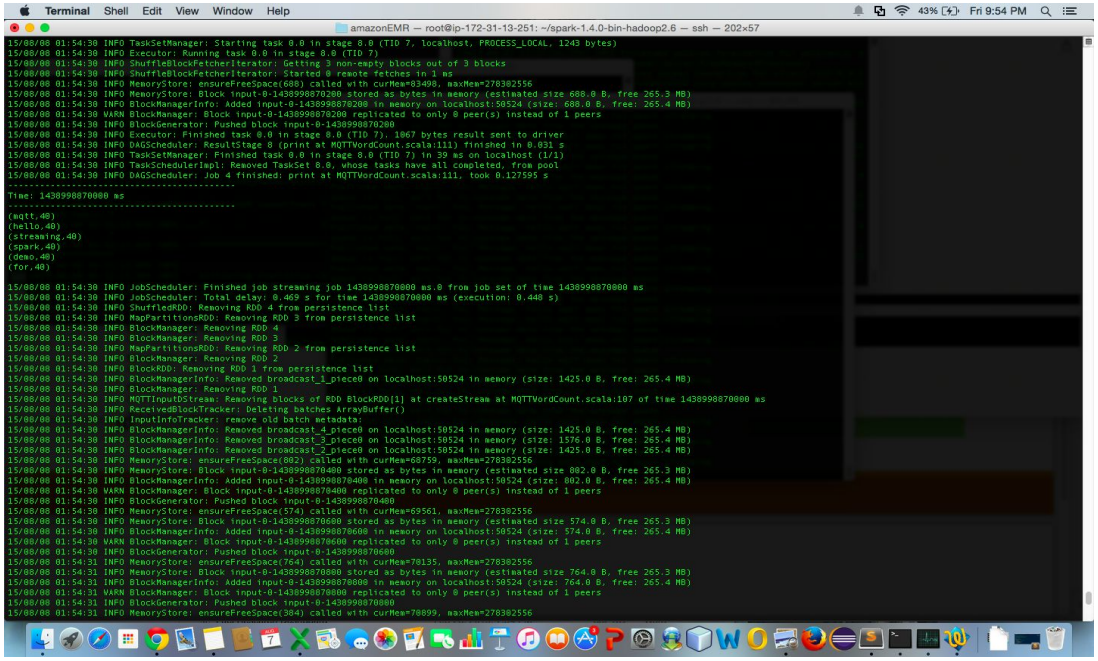

6. Steps for executing Twitter Popular Tags example (in Scala):

Twitter Credential Setup:

Since the whole exercise is based on twitter, it is necessary to configure authentication with a twitter account using a consumer key+secret pair and an access token+secret pair.

1. Open the twitter's application settings page (https://apps.twitter.com/) and create a new application by clicking on the "Create a new application" button and providing the required fields.

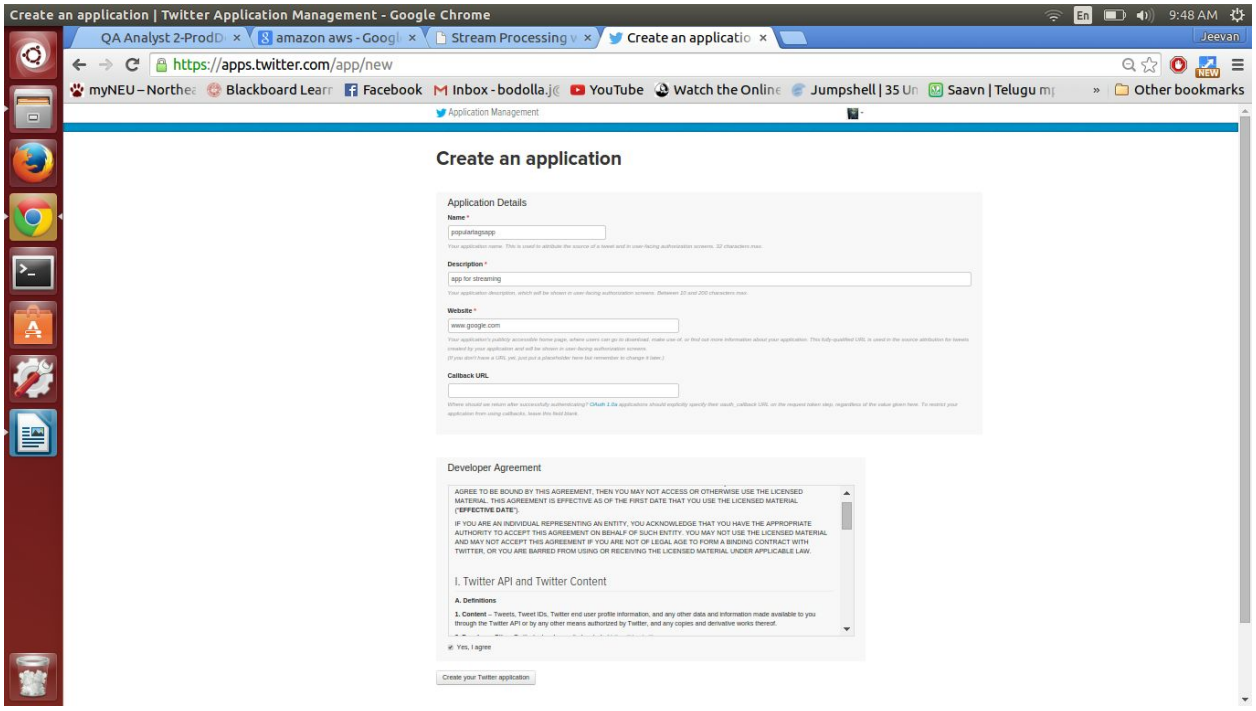

2. Once you create the application, click on the "Keys and Access Tokens" tab where you can see the API key and API secret that have been generated. Click on the "Create My Access Token" button at the bottom to generate the access token and access token secret.

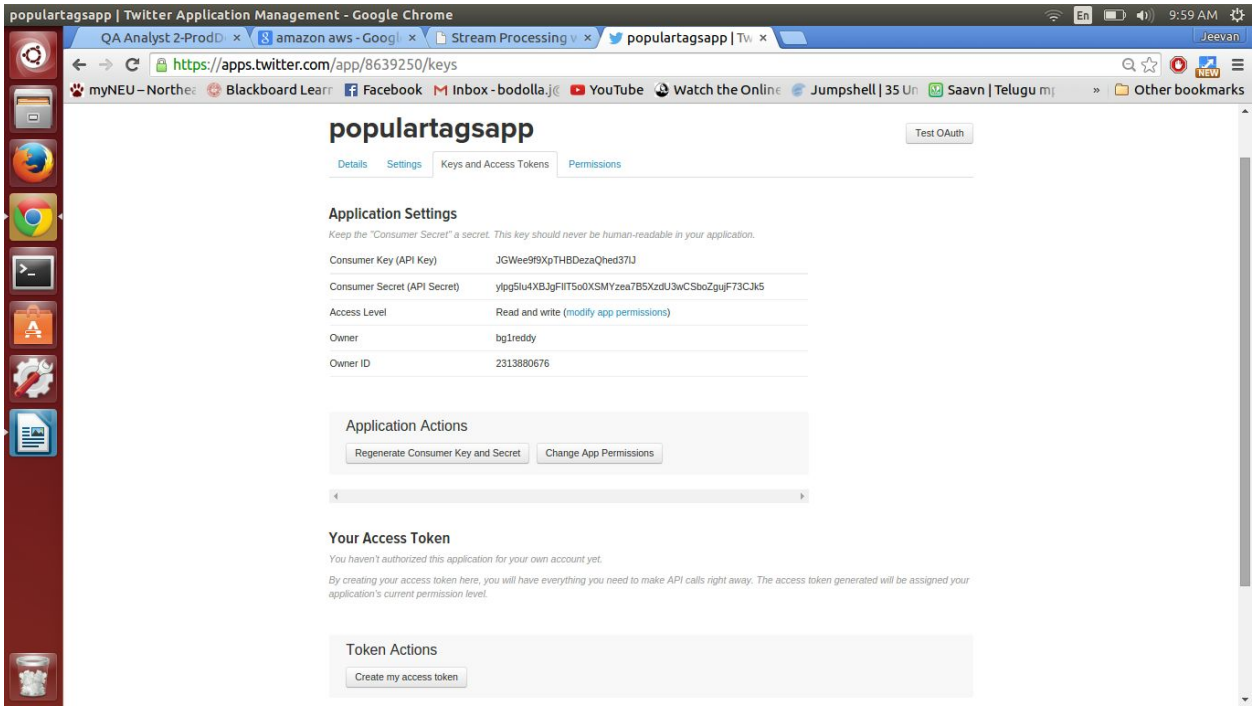

3. Finally, we will see the API Key, API Secret, Access token and Access token secret we are going to use in the next steps.

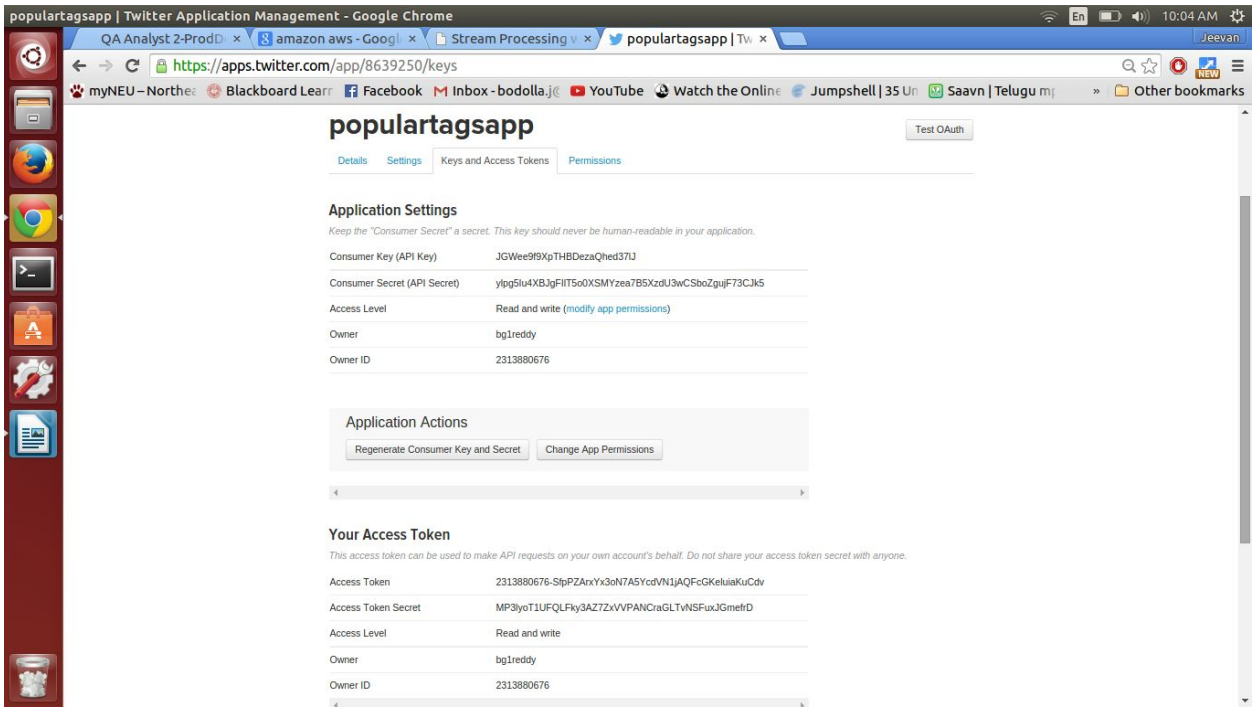

4. Now, SSH into the master node with the private key file used to launch the cluster by typing the following command:

$ ssh hadoop@ec2-52-1-104-201.compute-1.amazonaws.com -i ~/twitter.pem

5. Install Spark 1.4.0 on the EMR cluster by typing the following commands:

$ cd /usr/lib/
$ sudo wget
http://psg.mtu.edu/pub/apache/spark/spark-1.4.0/spark-1.4.0-bin-hadoop2.6.tgz
$ sudo tar -zxvf spark-1.4.0-bin-hadoop2.6.tgz
$ cd spark-1.4.0-bin-hadoop2.6

6. Finally, compile the TwitterPopularTags class with the API Key, API Secret, Access token and Access token secret as arguments as shown below:

$ bin/run-example org.apache.spark.examples.streaming.TwitterPopularTags
JeCdbwdDuVpPzwRh4aLvJYKj2
6cEXsnpvcaonemZz9KJG059E2yOXWLFCnQax0t1rdq9FRcPwIq
2313880676-QhO7bT3tojJfurFZJjjpJqAOXtj7fdB6gJPLzgk
XVTtZu0sYtd7WqBemaDf8bOXSEhUkqfUEwJDXv1MoRFJf

and you will soon find the popular tweets over sliding 10 and 60 second window from a twitter stream being printed on the screen something like this:

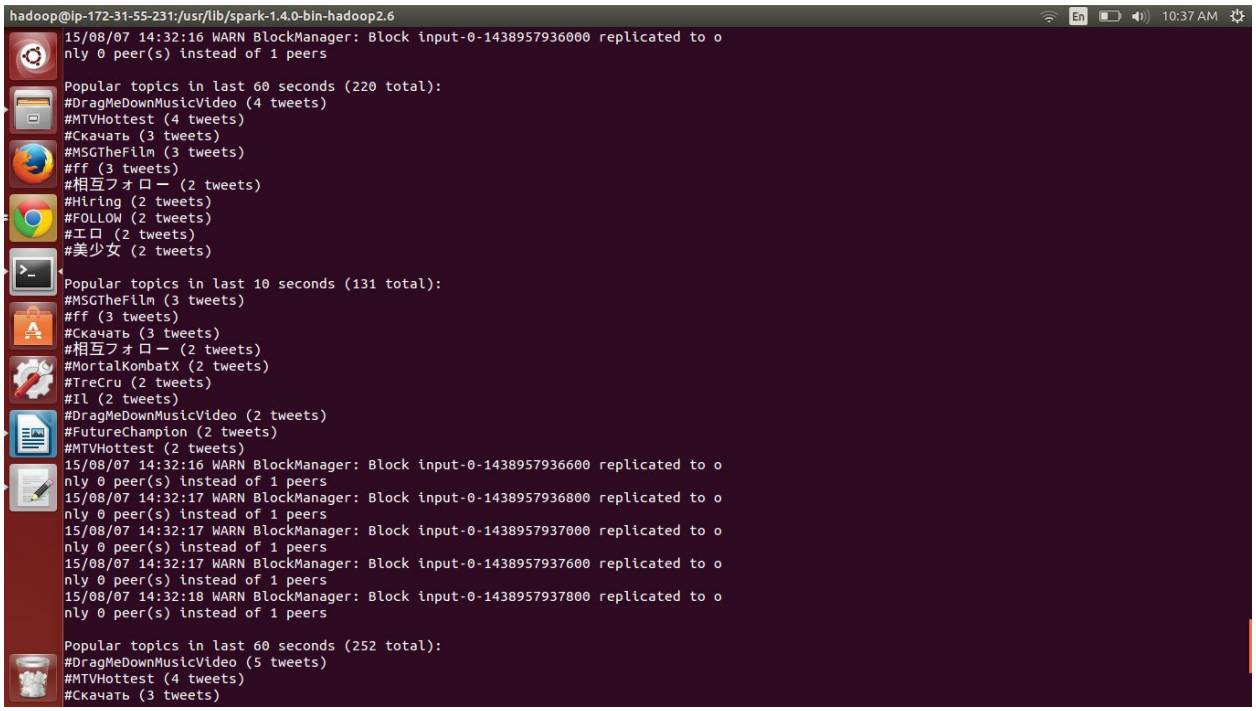

7. Steps for executing ZeroMQ wordcount example (in Scala)

- Install ZeroMQ by using the following command
  - $ apt-get install libzmq-dev

- Run the publisher by using the following command
  - $ bin/run-example org.apache.spark.examples.streaming.SimpleZeroMQPublisher tcp://127.0.1.1:1234 foo.ba

- Run the ZeroMQ example file as follows
  - $ bin/run-example org.apache.spark.examples.streaming.ZeroMQWordCount tcp://127.0.1.1:1234 foo

- The output will be something as below:

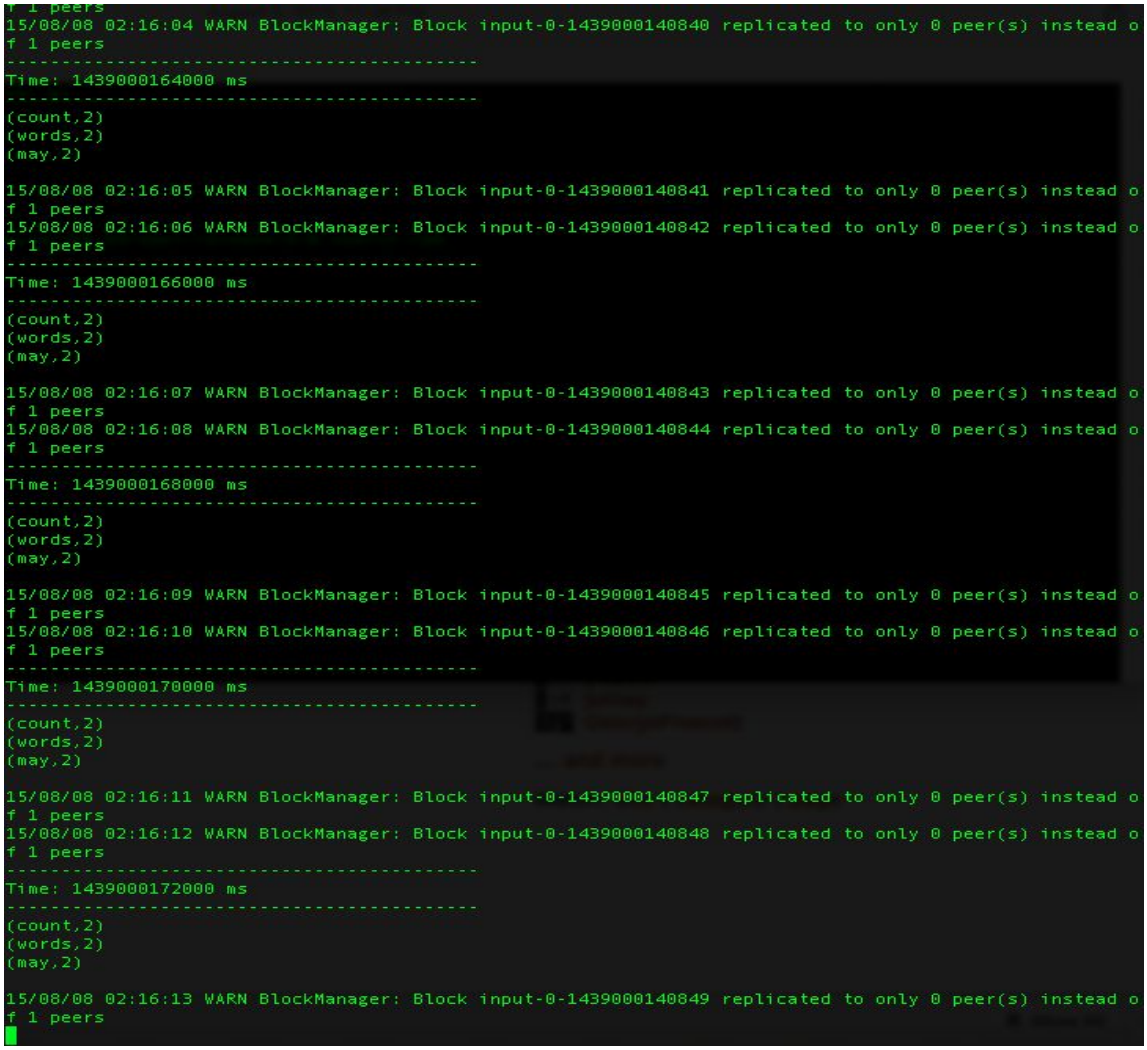

8. Steps for executing Kafka wordcount example (in Python)

- Prerequisite: You need to open at least **7** terminal windows connected to EMR
- Note- Kafka doesn't work on Spark 1.4.1 so we installed Spark 1.4.0

- EMR Master node doesn't have **wget** installed. so first we'll have to install that in order to be able to download any other stuff

    - $ sudo yum -y install wgetv
    - $ wget --no-check-certificate --no-cookies --header "Cookie: oraclelicense=accept-securebackup-cookie"http://download.oracle.com/otn-pub/java/jdk/7u67-b01/jdk-7u67-linux-x64.rpm

- Now we need to download Spark 1.4.0 as EMR comes with Spark 1.4.1
    - $ wget http://mirror.symnds.com/software/Apache/spark/spark-1.4.0/spark-1.4.0-bin-hadoop2.6.tgz
    - $ tar -xzf spark-1.4.0-bin-hadoop2.6.tgz

- Then we can download Kafka 2.11-0.8.2.1 from the mirror site using wget
    - $ wget http://apache.mirrors.tds.net/kafka/0.8.2.1/kafka_2.9.1-0.8.2.1.tgz
    - $ tar -xzf kafka_2.9.1-0.8.2.1.tgz
    - $ cd kafka_2.9.1-0.8.2.1

- Start Zookeeper by using the following command:

    $ bin/zookeeper-server-start.sh config/zookeeper.properties

    Zookeeper starts at localhost: 2181

- Start Kafka broker in another terminal by using the following command:

    $ cd kafka_2.9.1-0.8.2.1
    $ bin/kakfa-server-start.sh config/server.properties

    KafkaBroker starts at localhost: 9092

- Create a Kafka Topic of your choice in another terminal by using the following command:

$ cd kafka_2.9.1-0.8.2.1
$ bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1 --topic kafkatopic

This creates a topic by name Kafkatopic

- Now, start the producer in another terminal which produces data by using the following command:

$ cd kafka_2.9.1-0.8.2.1
$ bin/kafka-console-producer.sh --broker-list localhost:9092 --topic kafkatopic

Type in anything you would like to send to the kafka broker in the command window

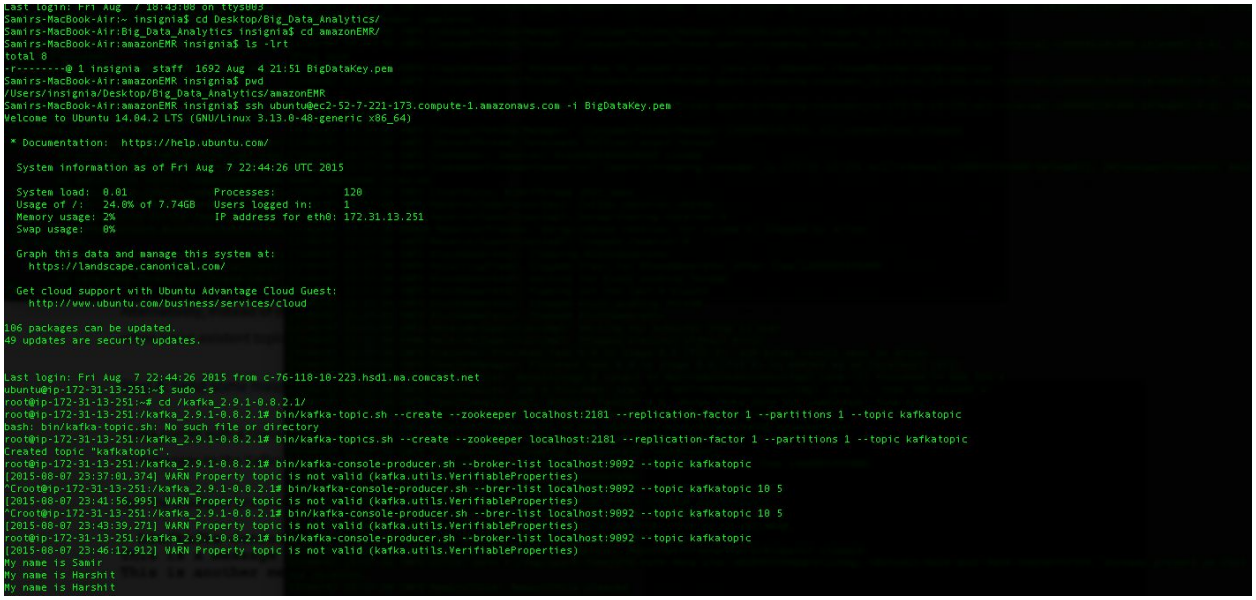

- Run the wordcount example in another terminal in spark 1.4.0 directory by using the following command:

$ bin/spark-submit --jars external/kafka-assembly/target/scala-*/*.jar examples/src/main/python/streaming/kafka_wordcount.py localhost:2181 kafkatopic

The output would be something like this:

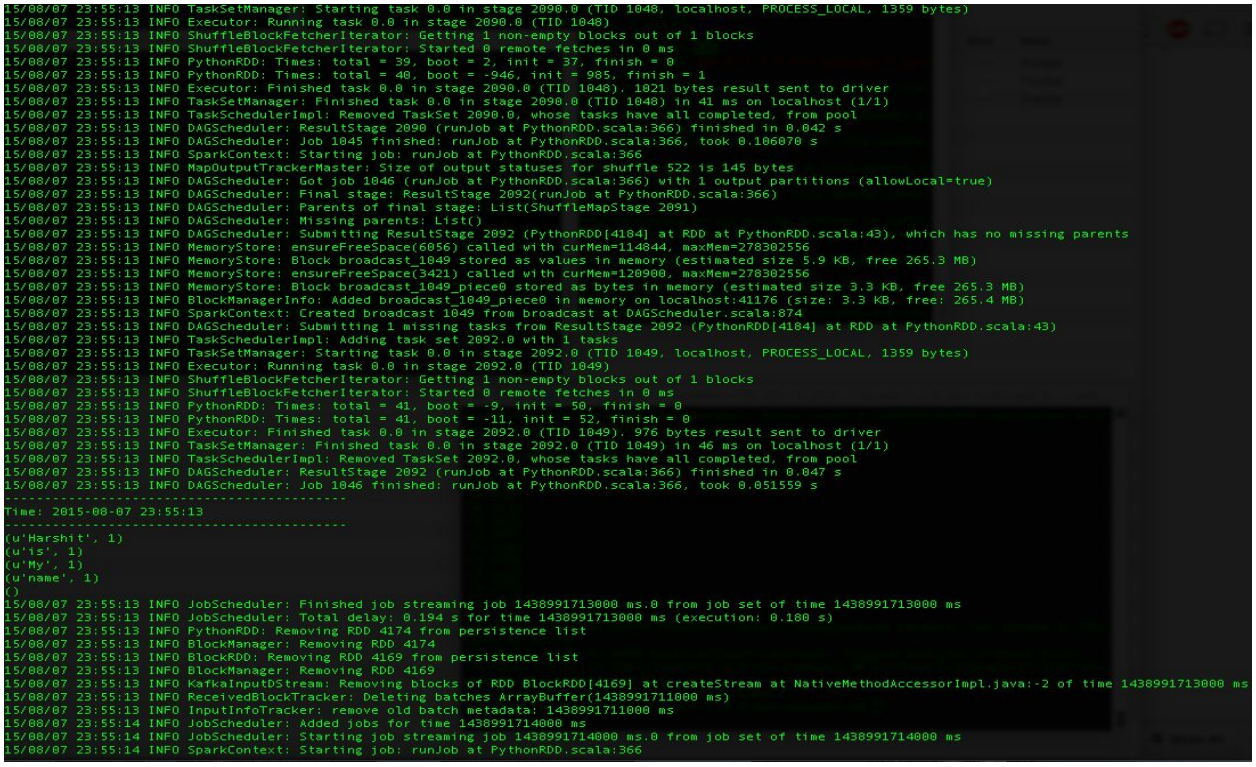

9. Steps for executing SQL Network Wordcount example (in Python)

- We need to start a Netcat server first in order to count the words over the network
    - $ nc -lk 9999

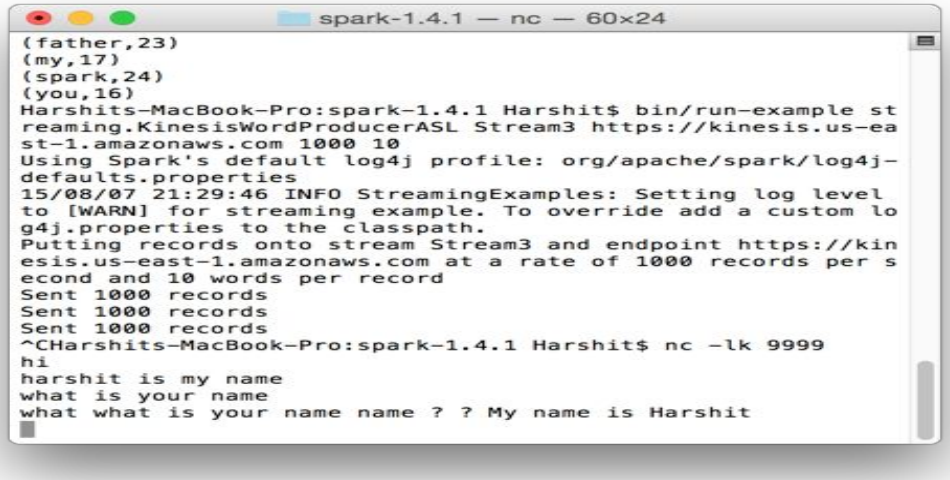

- Then we start a listener which count the number of words
    - $ bin/spark-submit examples/src/main/python/streaming/sql_network_wordcount.py localhost 9999

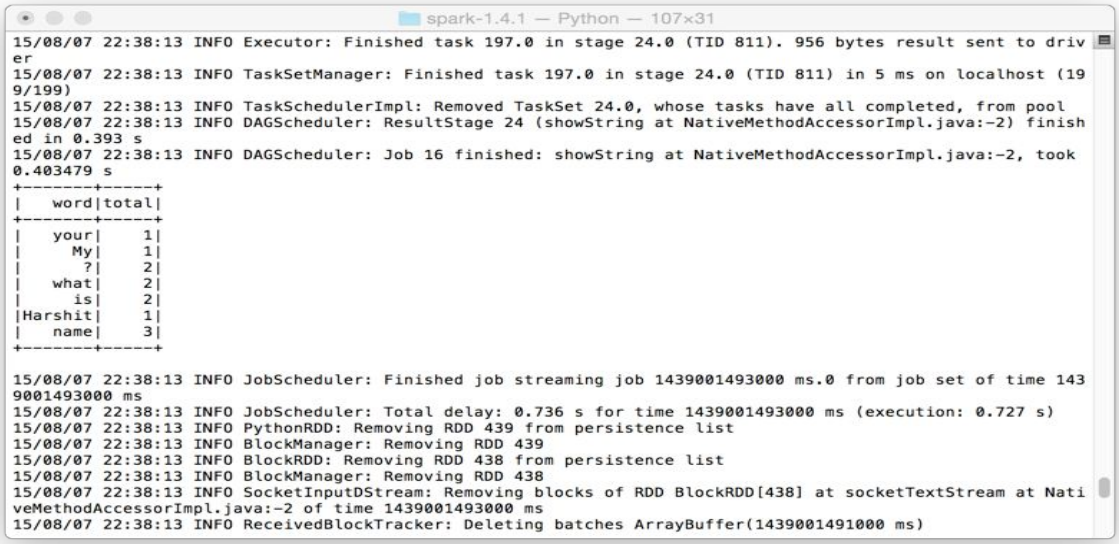

10. Steps for executing HDFS wordcount example (in Python)

- Change the directory to hadoop and create a folder called input in the hadoop folder as follows
    - $ cd /usr/lib/hadoop

- Make a directory in HDFS
    - $ bin/hadoop fs -mkdir /input

- Run the example from the spark directory with the hdfs directory as input
    - $ bin/spark-submit examples/src/main/python/streaming/hdfs_wordcount.py hdfs:///input

- Put a file into the hdfs directory for wordcount
    - $ bin/hadoop fs -put /usr/lib/spark/examples/src/main/resources/kv1.txt /input

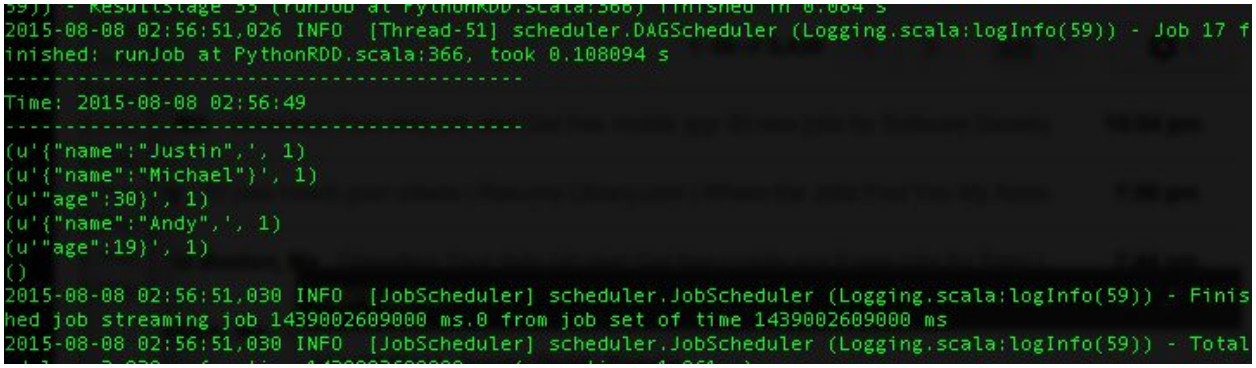
59)) - ResultStage 33 (runJob at PythonRDD.scala:366) finished in 0.084 s
2015-08-08 02:56:51,026 INFO [Thread-51] scheduler.DAGScheduler (Logging.scala:logInfo(59)) - Job 17 f
inished: runJob at PythonRDD.scala:366, took 0.108094 s
Time: 2015-08-08 02:56:49
(u'{"name":"Justin",', 1)
(u'{"name":"Michael"}', 1)
(u'"age":30}', 1)
(u'{"name":"Andy",', 1)
(u'"age":19}', 1)
()
2015-08-08 02:56:51,030 INFO [JobScheduler] scheduler.JobScheduler (Logging.scala:logInfo(59)) - Finis
hed job streaming job 1439002609000 ms.0 from job set of time 1439002609000 ms
2015-08-08 02:56:51,030 INFO [JobScheduler] scheduler.JobScheduler (Logging.scala:logInfo(59)) - Total